

Mikhail Kuznetsov

New York, NY mikkuzne@gmail.com +1 646 988 1199

Google Scholar LinkedIn

Summary

Research Scientist specializing in securing agentic AI systems and foundation models for audit/security logs. Current focus: detecting intent deviation and policy violations in agent runtime traces, extending detections with OS-level signals (eBPF), auto-detecting agentic workloads in dynamic environments, and co-evolutionary red/blue-team auditing of multi-agent systems. Led an audit-log FM end-to-end: diversity-preserving 10B $\rightarrow$ 10M training-data sampling, pretraining + contrastive fine-tuning, interpretable anomaly scoring, and evaluation with limited feedback. Ongoing research in LLM alignment and robust self-supervised learning.

Research Interests

Agentic AI security • Automated red-teaming & co-evolutionary auditing • Foundation models for logs • Robust/self-supervised learning • Anomaly detection • Representation learning for structured/tabular data • Evaluation & interpretability • LLM alignment

Experience

Amazon AWS – AI Security and Observability

2022 – Present

Senior Applied Scientist

- **Agentic AI Security:** Leading detection for securing AI agents in production. Models over agent runtime traces (e.g. Amazon Bedrock AgentCore) detect intent deviation and policy violations (prompt injection being one driver), conditioned on an agent’s historical behavior and agentic context (user/system prompts). Extending detections with OS-level telemetry (eBPF, system logs) to verify what an agent actually did against what its trace claims, and to auto-detection of agentic workloads in dynamic environments.
- **Kernel-Level Evidence for Agent Security:** Co-authored the first paired-evidence study of kernel-level vs. application-layer signal for agent security. Built ACE, a corpus of 4,047 agent sessions pairing syscall traces with tool manifests and agent transcripts across 17 threat models (14 of 25 OWASP LLM/agentic categories). Across 10 detectors in four families, kernel evidence is discriminative on its own, cross-layer composition generally beats either single layer, and detectors generalize to unseen attack families and transfer to another agent runtime.
- **Co-Evolutionary Red-Teaming of Multi-Agent Systems:** Co-developed PurpleAudit (NeurIPS 2026, Evaluations & Datasets), an auditing framework where red and blue teams co-evolve over execution traces, generating progressively stronger attacks and the defenses that answer them. Targets *task hijacking*: injections that redirect an MAS toward an attacker-chosen task inside its own domain, with no malicious verb to refuse. Reached 31.26% attack success on a frontier backend where all 22 baseline attacks fail; the co-evolved prompt-level defense holds at negligible utility cost and transfers across backends.
- **LLM-Based Foundation Model on Audit Logs:** Tech lead of a foundation model for audit logs, learning embeddings for security entities (IP, API, usernames) from dynamic audit data. The model was deployed to production in Sep 2025, enhancing Customer-facing anomaly detection, reducing False Positives by 20–30%. Main contributions are:

Training data creation. Designed a sampling algorithm that compresses $\sim$ 10B raw log events to $\sim$ 10M high-value examples while preserving entity diversity; runs on a regular cadence to continually refresh training data.

Contrastive Fine-tuning. Introduced contrastive fine-tuning for tenant-specific contexts, which are focused on entities particularly important for downstream consumers (asnOrg, api).

Evaluation & interpretability. Built a multi-signal evaluation framework (weak supervision, third-party feeds, expert labels) and interpretability tooling (embedding similarity, contextual neighbors) for explainable anomalies.

Architectural design (hierarchical logs). Co-designed a memory-efficient architecture leveraging the nested structure of log entries for long-context, high-throughput security workload.

GenAI in Security Benchmarking: Co-authored a benchmark evaluating generative models in security workflows.

- **Robust/Secure Learning:** Co-developed distributionally robust self-supervised learning for tabular data; advanced graph-based domain reputation with weak supervision; ongoing research on agentic AI security for context-aware anomaly/failure detection in multi-agent systems.

Yahoo! Research, New York
Research Scientist

2017 – 2021

- **ExtremeText:** contributed to a no-regret generalization of hierarchical softmax for *extreme* multi-label classification; improved training/inference efficiency and accuracy over very large label spaces.
- **VisualTextRank:** designed an unsupervised, graph-based content extraction method to automate ad text → image matching; formulated a TextRank-style ranking over multimodal (text/vision) features, built data/eval pipelines, and improved retrieval relevance in internal benchmarks.
- Built unified multi-task CTR/CVR models that improved advertising ROI by >10% in online A/B tests.

Education

Moscow Institute of Physics and Technology

Ph.D. (Computer Science, Mathematical Modeling), 2016

M.S. with Honors, Applied Mathematics & Physics, 2013

B.S. with Honors, Applied Mathematics & Physics, 2011

Yandex School of Data Analysis: Diploma in Data Science, 2012

Publications by Topic

AI Security & Observability

- *PurpleAudit: A Co-Evolutionary Auditing Framework for Task Hijacking in Multi-Agent Systems.* T. Bilot, Z. Zhang, K. Liu, **M. Kuznetsov**, W. Ding. NeurIPS (Evaluations & Datasets Track), 2026.
- *On the Effectiveness of Kernel-Level Evidence for Agent Security.* S. King, Z. Zhang, **M. Kuznetsov**, K. Liu, B. Coskun, W. Ding. arXiv:2609.28915, in submission, 2026.
- *HLogformer: A Hierarchical Transformer for Representing Log Data.* Z. Hou, M. Ghashami, **M. Kuznetsov**, MA. Torkamani. 2024
- *An Evaluation Benchmark for Generative AI in Security Domain.* M. Ghashami, **M. Kuznetsov**, V. Gao, G. Teng, P. Wallis, J. Xie, A. Torkamani, B. Coskun, W. Ding. KDD GenAI Evaluation Workshop, 2024.

Robust / Extreme Classification

- *Distributionally Robust Self-Supervised Learning for Tabular Data.* S. Ghosh, J. Xie, **M. Kuznetsov**. NeurIPS Third Table Representation Learning Workshop, 2024.
- *A No-Regret Generalization of Hierarchical Softmax to Extreme Multi-Label Classification.* M. Wydmuch, K. Jasinska, **M. Kuznetsov**, R. Busa-Fekete, K. Dembczynski. NeurIPS, 2018.

Language Model Alignment

- *STARS: Synchronous Token Alignment for Robust Supervision in Large Language Models.* M. Quamar,

M. Areeb, **M. Kuznetsov**, M. Ozmen, B. Celik. SPIGM Workshop, ICML, 2026.

- *Adaptive Blockwise Search: Inference-Time Alignment for Large Language Models*. M. Quamar, M. Areeb, N. Sharma, Nishant, A. Shreekumar, J. Jonathan, M. Ozmen, **M. Kuznetsov**, B. Celik. 2025.

Generative Models & Retrieval for Ads/Vision

- *Salient Object-Aware Background Generation using Text-Guided Diffusion Models*. A. Eshratifar, J. Soares, K. Thadani, S. Mishra, **M. Kuznetsov**, Y. Ku, P. De Juan. CVPR Workshop on Generative Models for Computer Vision, 2024.
- *Staging E-Commerce Products for Online Advertising using Retrieval Assisted Image Generation*. Y. N. Ku, **M. Kuznetsov**, S. Mishra, P. de Juan. AdKDD, 2023.
- *VisualTextRank: Unsupervised Graph-Based Content Extraction for Automating Ad Text-to-Image Search*. KDD (ADS Track). S. Mishra, **M. Kuznetsov**, G. Srivastava, M. Sviridenko, 2021.

Open-Source Projects

ClawGuard: AI workload observer that auto-detects and watches agentic workloads on a machine (processes, network, parent chains) via an LLM loop maintaining a live world model in place of hand-written rules. github.com/mikkuzne/clawguard

apparty: Telegram bot that builds sandboxed (**bwrap**) web apps on demand with an Anthropic tool-use agent; a second model narrates the agent's actions in plain English. github.com/mikkuzne/apparty

pitchclaw: LLM-curated calibrated priors for football match outcomes, built on a weekly-rewritten team-strength model feeding a mechanical decision filter. github.com/mikkuzne/pitchclaw

Skills

Data & Infra: PySpark; Pandas; Arrow/Parquet; large-scale data curation/sampling; AWS (SageMaker, S3, EMR, EKS).

Modeling: scikit-learn; PyTorch; Hugging Face (Transformers, Accelerate, Datasets, PEFT); vLLM; LoRA/QLoRA.